

Longest Increasing Subsequence

Two classic approaches - pick based on constraints:

1) $O(n^2)$ DP:

Let $dp[i]$ = length of the LIS ending at i .

Transition:

- For each i , check all $j < i$:

- If $nums[j] < nums[i]$, we can extend: $dp[i] = \max(dp[i], dp[j] + 1)$.

- Initialize all $dp[i] = 1$.

- Answer = $\max_i dp[i]$.

Why it works: you extend the best increasing subsequence that ends before i and is smaller than $nums[i]$.

2) $O(n \log n)$ "patience sorting" trick (preferred for larger n):

Maintain an array $tails$, where:

- $tails[k]$ = the smallest possible tail value of an increasing subsequence of length $k+1$.

Process each x in $nums$:

- Find the first index i in $tails$ with $tails[i] \geq x$ (lower bound).

- If found, set $tails[i] = x$ (we improved the tail for length $i+1$).

- If not found (i.e. x is larger than all tails), append x to $tails$ (we extended the LIS by 1).

- The length of LIS is $tails.size()$ at the end.

Key intuition:

- $tails$ is kept sorted.

- We never claim $tails$ itself is a valid subsequence; it's a bookkeeping structure ensuring minimal tails for each length.

- Using $>=$ gives strictly increasing subsequences (duplicates don't extend).

Complexity:

- Time: $O(n \log n)$ for the binary searches.

- Space: $O(n)$ for $tails$.