

Word Ladder

This problem can be solved efficiently using Breadth-First Search (BFS).

Key Concepts:

1. Graph Construction:

- Words are nodes.
- An edge exists between two words if they differ by exactly one letter.
- Example: hot \leftrightarrow dot, hot \leftrightarrow lot.

2. Shortest Path Problem:

- You are finding the shortest path from begin word to end word in this implicit graph.
- BFS is ideal since all edges have equal weight.

3. Constraints to Check:

- If end word is not in word list, the answer is 0.
- You can transform only one character at a time.
- You must only use words from word list except for begin word.

BFS Strategy:

1. Preprocess:

- Use a set for word list for fast lookup.

2. BFS Queue:

- Each element is a (word, depth) pair.
- depth represents how many transformations taken so far (starting at 1 with begin word).

3. At each BFS level:

- Try changing each character in the word to all 26 letters.
- If the new word is in the word list and not visited:
 - Add it to the queue and mark it as visited.
 - If it matches end word, return depth + 1.

4. If queue gets exhausted: return 0.

Optimization ideas:

- You can preprocess intermediate states like list, hash to reduce search space.
- Even better: use Bidirectional BFS to cut the search space in half (not required for correctness, but improves performance).

Complexity:

- Time: $O(N * M^2)$, where N is number of words, M is word length.
- Space: $O(N * M)$ for the queue and visited set